

FDE Project Evaluation



Barath Velmurugan

← Re-Evaluate

Evaluation Results Summary

Evaluation Complete!

Input: <https://github.com/barathvelmu/fde-enterprise-agent>

TOTAL SCORE

95 / 95

Breakdown by Criteria

Tool & Interface Design

20 / 20 pts

Reason: The code excellently fulfills all criteria by providing detailed tool docstrings, using highly specific tool names like 'create_critical_incident_ticket', enforcing explicit Pydantic JSON schemas for validation, and implementing robust error handling that returns clear, actionable recovery instructions to the LLM.

Context & Memory

20 / 20 pts

Reason: The code excellently implements all four criteria: a comprehensive system constitution prompt defines the agent's persona and constraints, history compaction is achieved through token-budgeted sliding windows, persistent session state is managed via mocked Vertex AI datastores, and memory consolidation is handled as non-blocking async tasks.

Orchestration & Logic

20 / 20 pts

Reason: The codebase successfully implements all criteria, utilizing ADK Coordinator and Sequential agents, strategic model routing between Gemini Flash and Pro, policy check guardrails, and explicit human-in-the-loop approval for high-severity actions.

Observability & Tracing

20 / 20 pts

Reason: The code excellently implements all observability criteria: it uses Python's logging module to output structured JSON payloads, explicitly records pre-execution intentions and post-execution outcomes in the ticket creation tool, utilizes OpenTelemetry spans for distributed tracing, and applies active PII redaction (via regex scrubbing) before logging and memory consolidation.

Infrastructure & CI/CD

15 / 15 pts

Reason: The repository successfully fulfills all criteria: it includes an automated evaluation suite against a golden dataset, documents Terraform IaC configurations and agent CLI deployment commands, and utilizes Google Cloud Secret Manager to securely manage secrets without any hardcoded API keys.

Download Results CSV